The counties dataset

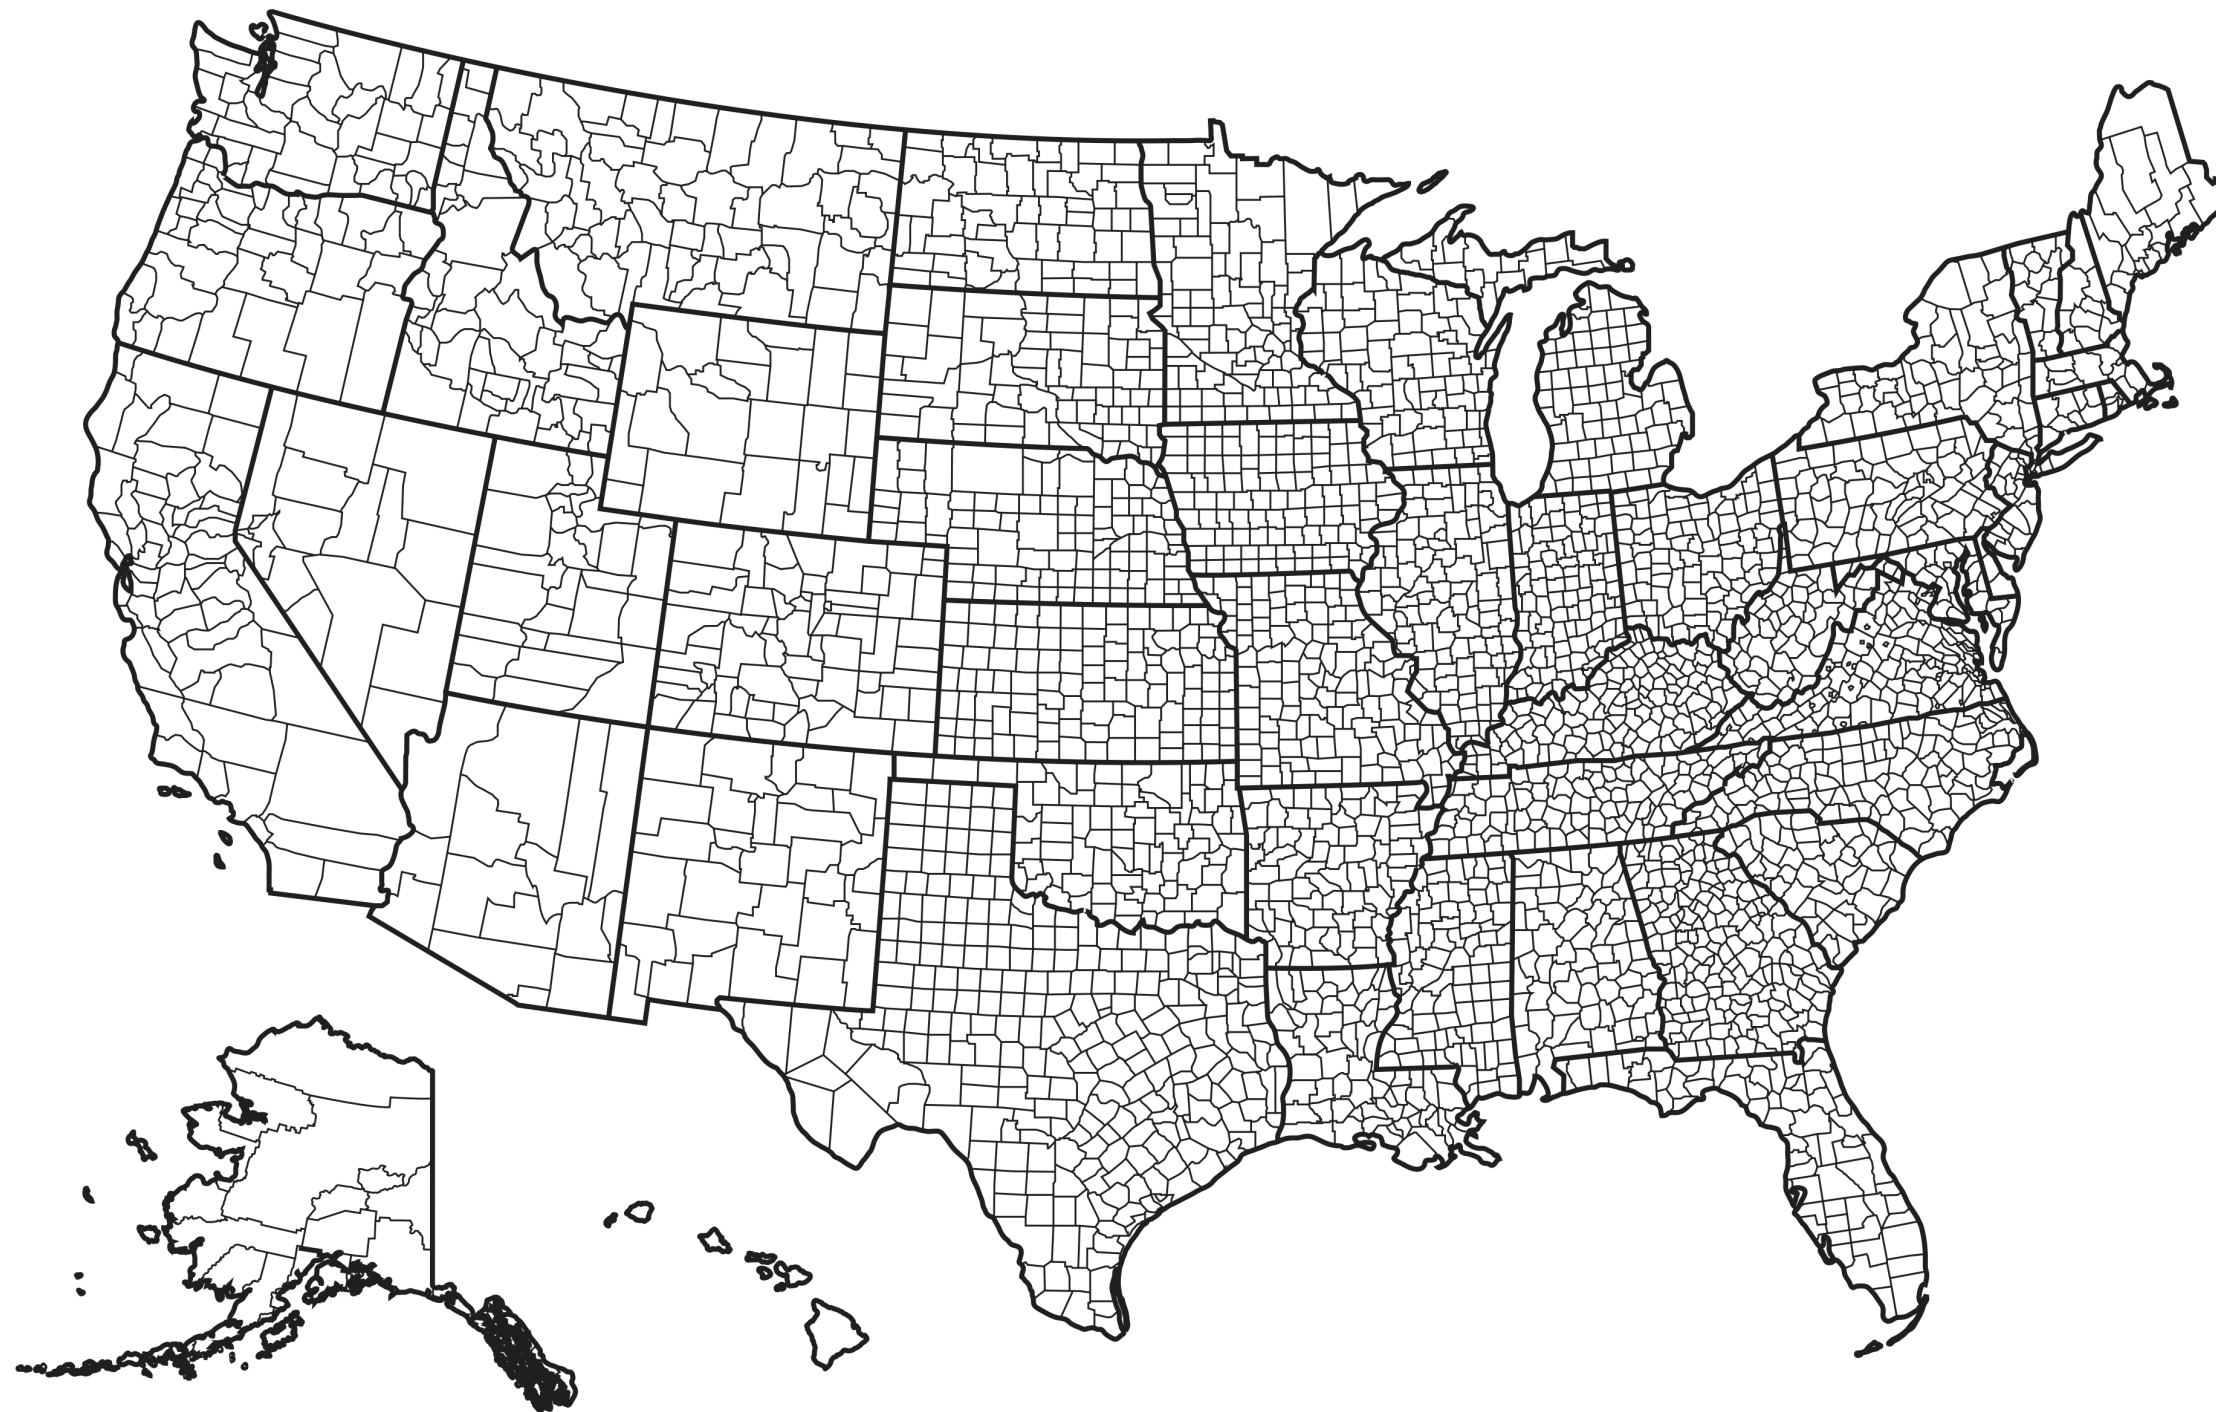## DATA MANIPULATION WITH DPLYR

Chris Cardillo
Data Scientist

Chapter 1 verbs

- `select()`

- `filter()`

- `arrange()`

- `mutate()`

2015 United States Census

Counties dataset

counties

```
# A tibble: 3,138 x 40
   census_id state county region metro population    men  women hispanic white black native asian pacific
   <chr>     <chr> <chr>  <chr>  <chr>      <dbl>   <dbl>  <dbl>    <dbl> <dbl> <dbl>  <dbl> <dbl>   <dbl>
 1 1001      Alab… Autau… South  Metro      55221   26745  28476      2.6  75.8  18.5    0.4   1         0
 2 1003      Alab… Baldw… South  Metro     195121   95314  99807      4.5  83.1   9.5    0.6   0.7       0
 3 1005      Alab… Barbo… South  Nonm…      26932   14497  12435      4.6  46.2  46.7    0.2   0.4       0
 4 1007      Alab… Bibb   South  Metro      22604   12073  10531      2.2  74.5  21.4    0.4   0.1       0
 5 1009      Alab… Blount South  Metro      57710   28512  29198      8.6  87.9   1.5    0.3   0.1       0
 6 1011      Alab… Bullo… South  Nonm…      10678    5660   5018      4.4  22.2  70.7    1.2   0.2       0
 7 1013      Alab… Butler South  Nonm…      20354    9502  10852      1.2  53.3  43.8    0.1   0.4       0
 8 1015      Alab… Calho… South  Metro     116648   56274  60374      3.5  73    20.3    0.2   0.9       0
 9 1017      Alab… Chamb… South  Nonm…      34079   16258  17821      0.4  57.3  40.3    0.2   0.8       0
10 1019      Alab… Chero… South  Nonm…      26008   12975  13033      1.5  91.7   4.8    0.6   0.3       0
# … with 3,128 more rows, and 26 more variables: citizens <dbl>, income <dbl>, income_err <dbl>,
#   income_per_cap <dbl>, income_per_cap_err <dbl>, poverty <dbl>, child_poverty <dbl>,
#   professional <dbl>, service <dbl>, office <dbl>, construction <dbl>, production <dbl>, drive <dbl>,
#   carpool <dbl>, transit <dbl>, walk <dbl>, other_transp <dbl>, work_at_home <dbl>, mean_commute <dbl>,
#   employed <dbl>, private_work <dbl>, public_work <dbl>, self_employed <dbl>, family_work <dbl>,
#   unemployment <dbl>, land_area <dbl>
```

```
glimpse(counties)
```

```
Observations: 3,138
Variables: 40
$ census_id      <chr> "1001", "1003", "1005", "1007", "1009", "1011", "1013", …
$ state          <chr> "Alabama", "Alabama", "Alabama", "Alabama", "Alabama", "…
$ county         <chr> "Autauga", "Baldwin", "Barbour", "Bibb", "Blount", "Bull…
$ region         <chr> "South", "South", "South", "South", "South", "South", "S…
$ metro          <chr> "Metro", "Metro", "Nonmetro", "Metro", "Metro", "Nonmetr…
$ population     <dbl> 55221, 195121, 26932, 22604, 57710, 10678, 20354, 116648…
$ men            <dbl> 26745, 95314, 14497, 12073, 28512, 5660, 9502, 56274, 16…
$ women          <dbl> 28476, 99807, 12435, 10531, 29198, 5018, 10852, 60374, 1…
$ hispanic       <dbl> 2.6, 4.5, 4.6, 2.2, 8.6, 4.4, 1.2, 3.5, 0.4, 1.5, 7.6, 0…
$ white          <dbl> 75.8, 83.1, 46.2, 74.5, 87.9, 22.2, 53.3, 73.0, 57.3, 91…
$ black          <dbl> 18.5, 9.5, 46.7, 21.4, 1.5, 70.7, 43.8, 20.3, 40.3, 4.8,…
$ native         <dbl> 0.4, 0.6, 0.2, 0.4, 0.3, 1.2, 0.1, 0.2, 0.2, 0.6, 0.4, 0…
$ asian          <dbl> 1.0, 0.7, 0.4, 0.1, 0.1, 0.2, 0.4, 0.9, 0.8, 0.3, 0.3, 0…
$ pacific        <dbl> 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0…
$ citizens       <dbl> 40725, 147695, 20714, 17495, 42345, 8057, 15581, 88612, …
$ income         <dbl> 51281, 50254, 32964, 38678, 45813, 31938, 32229, 41703, …
...
```

Select

```
counties %>%
  select(state, county, population, unemployment)
```

```
# A tibble: 3,138 x 4
   state   county   population unemployment
   <chr>   <chr>         <dbl>        <dbl>
 1 Alabama Autauga       55221          7.6
 2 Alabama Baldwin      195121          7.5
 3 Alabama Barbour       26932         17.6
 4 Alabama Bibb          22604          8.3
 5 Alabama Blount        57710          7.7
 6 Alabama Bullock       10678         18
 7 Alabama Butler        20354         10.9
 8 Alabama Calhoun      116648         12.3
 9 Alabama Chambers      34079          8.9
10 Alabama Cherokee      26008          7.9
# … with 3,128 more rows
```

Creating a new table

```
counties_selected <- counties %>%
  select(state, county, population, unemployment)
```

```
# A tibble: 3,138 x 4
   state   county      population unemployment
   <chr>   <chr>            <dbl>        <dbl>
 1 Alabama Autauga          55221          7.6
 2 Alabama Baldwin         195121          7.5
 3 Alabama Barbour          26932         17.6
 4 Alabama Bibb             22604          8.3
 5 Alabama Blount           57710          7.7
 6 Alabama Bullock          10678         18
 7 Alabama Butler           20354         10.9
 8 Alabama Calhoun         116648         12.3
 9 Alabama Chambers         34079          8.9
10 Alabama Cherokee         26008          7.9
# … with 3,128 more rows
```

Let's practice!

DATA MANIPULATION WITH DPLYR

The filter and arrange verbs

DATA MANIPULATION WITH DPLYR

Chris Cardillo
Data Scientist

```
counties_selected <- counties %>%
  select(state, county, population, unemployment)


counties_selected
```

```
# A tibble: 3,138 x 4
   state   county    population unemployment
   <chr>   <chr>          <dbl>        <dbl>
 1 Alabama Autauga        55221          7.6
 2 Alabama Baldwin       195121          7.5
 3 Alabama Barbour        26932         17.6
 4 Alabama Bibb           22604          8.3
 5 Alabama Blount         57710          7.7
 6 Alabama Bullock        10678         18
 7 Alabama Butler         20354         10.9
 8 Alabama Calhoun       116648         12.3
 9 Alabama Chambers       34079          8.9
10 Alabama Cherokee       26008          7.9
# … with 3,128 more rows
```

Arrange

```
counties_selected %>%
  arrange(population)
```

```
# A tibble: 3,138 x 4
   state       county     population unemployment
   <chr>       <chr>          <dbl>        <dbl>
 1 Hawaii      Kalawao           85          0
 2 Texas       King             267          5.1
 3 Nebraska    McPherson        433          0.9
 4 Montana     Petroleum        443          6.6
 5 Nebraska    Arthur           448          4
 6 Nebraska    Loup             548          0.7
 7 Nebraska    Blaine           551          0.7
 8 New Mexico  Harding          565          6
 9 Texas       Kenedy           565          0
10 Colorado    San Juan         606         13.8
# … with 3,128 more rows
```

Arrange: descending

```
counties_selected %>%
  arrange(desc(population))
```

```
# A tibble: 3,138 x 4
   state      county       population unemployment
   <chr>      <chr>             <dbl>        <dbl>
 1 California Los Angeles    10038388         10
 2 Illinois   Cook            5236393         10.7
 3 Texas      Harris          4356362          7.5
 4 Arizona    Maricopa        4018143          7.7
 5 California San Diego       3223096          8.7
 6 California Orange          3116069          7.6
 7 Florida    Miami-Dade      2639042         10
 8 New York   Kings           2595259         10
 9 Texas      Dallas          2485003          7.6
10 New York   Queens          2301139          8.6
# … with 3,128 more rows
```

Filter

```
counties_selected %>%
  arrange(desc(population)) %>%
  filter(state == "New York")
```

```
# A tibble: 62 x 4
   state    county      population unemployment
   <chr>    <chr>            <dbl>        <dbl>
 1 New York Kings          2595259           10
 2 New York Queens         2301139          8.6
 3 New York New York       1629507          7.5
 4 New York Suffolk        1501373          6.4
 5 New York Bronx          1428357           14
 6 New York Nassau         1354612          6.4
 7 New York Westchester     967315          7.6
 8 New York Erie            921584            7
 9 New York Monroe          749356          7.7
10 New York Richmond        472481          6.9
# … with 52 more rows
```

Filter

```
counties_selected %>%
  arrange(desc(population)) %>%
  filter(unemployment < 6)
```

```
# A tibble: 949 x 4
   state    county       population unemployment
   <chr>    <chr>             <dbl>        <dbl>
 1 Virginia Fairfax         1128722          4.9
 2 Utah     Salt Lake       1078958          5.8
 3 Hawaii   Honolulu         984178          5.6
 4 Texas    Collin           862215          4.9
 5 Texas    Denton           731851          5.7
 6 Texas    Fort Bend        658331          5.1
 7 Kansas   Johnson          566814          4.5
 8 Maryland Anne Arundel     555280          5.9
 9 Colorado Jefferson        552344          5.9
10 Utah     Utah             551957          5.5
# … with 939 more rows
```

Combining conditions

```
counties_selected %>%
  arrange(desc(population)) %>%
  filter(state == "New York",
          unemployment < 6)
```

```
# A tibble: 5 x 4
  state    county     population unemployment
  <chr>    <chr>           <dbl>        <dbl>
1 New York Tompkins       103855          5.9
2 New York Chemung         88267          5.4
3 New York Madison         72427          5.1
4 New York Livingston      64801          5.4
5 New York Seneca          35144          5.5
```

Let's practice!

DATA MANIPULATION WITH DPLYR

Mutate

DATA MANIPULATION WITH DPLYR

Chris Cardillo
Data Scientist

```
counties_selected <- counties %>%
  select(state, county, population, unemployment)


counties_selected
```

```
# A tibble: 3,138 x 4
   state   county    population unemployment
   <chr>   <chr>          <dbl>        <dbl>
 1 Alabama Autauga        55221          7.6
 2 Alabama Baldwin       195121          7.5
 3 Alabama Barbour        26932         17.6
 4 Alabama Bibb           22604          8.3
 5 Alabama Blount         57710          7.7
 6 Alabama Bullock        10678         18
 7 Alabama Butler         20354         10.9
 8 Alabama Calhoun       116648         12.3
 9 Alabama Chambers       34079          8.9
10 Alabama Cherokee       26008          7.9
# … with 3,128 more rows
```

Total number of unemployed people

`population * unemployment / 100`

Mutate

```
counties_selected %>%
  mutate(unemployed_population = population * unemployment / 100)
```

```
# A tibble: 3,138 x 5
   state   county    population unemployment unemployed_population
   <chr>   <chr>          <dbl>        <dbl>                 <dbl>
 1 Alabama Autauga        55221          7.6                 4197.
 2 Alabama Baldwin       195121          7.5                14634.
 3 Alabama Barbour        26932         17.6                 4740.
 4 Alabama Bibb           22604          8.3                 1876.
 5 Alabama Blount         57710          7.7                 4444.
 6 Alabama Bullock        10678         18                   1922.
 7 Alabama Butler         20354         10.9                 2219.
 8 Alabama Calhoun       116648         12.3                14348.
 9 Alabama Chambers       34079          8.9                 3033.
10 Alabama Cherokee       26008          7.9                 2055.
# … with 3,128 more rows
```

```
counties_selected %>%
  mutate(unemployed_population = population * unemployment / 100) %>%
  arrange(desc(unemployed_population))
```

```
# A tibble: 3,138 x 5
   state      county          population unemployment unemployed_population
   <chr>      <chr>                <dbl>        <dbl>                 <dbl>
 1 California Los Angeles       10038388           10              1003839.
 2 Illinois   Cook               5236393         10.7               560294.
 3 Texas      Harris             4356362          7.5               326727.
 4 Arizona    Maricopa           4018143          7.7               309397.
 5 California Riverside          2298032         12.9               296446.
 6 California San Diego          3223096          8.7               280409.
 7 Michigan   Wayne              1778969         14.9               265066.
 8 California San Bernardino     2094769         12.6               263941.
 9 Florida    Miami-Dade         2639042           10               263904.
10 New York   Kings              2595259           10               259526.
# … with 3,128 more rows
```

Let's practice!

DATA MANIPULATION WITH DPLYR